

Searching

Problem: Search

- We are given a list of records.
- Each record has an associated key.
- Give efficient algorithm for searching for a record containing a particular key.
- Efficiency is quantified in terms of average time analysis (number of comparisons) to retrieve an item.

Search

[0]


[1]

[2]

[3]

[4]

[700]



Each record in list has an associated key.
In this example, the keys are ID numbers.

Given a particular key, how can we
efficiently retrieve the record from the list?



Serial Search

- Step through array of records, one at a time.
- Look for record with matching key.
- Search stops when
 - record with matching key is found
 - or when search has examined all records without success.

Pseudocode for Serial Search

```
// Search for a desired item in the n array elements  
// starting at a[first].  
// Returns pointer to desired record if found.  
// Otherwise, return NULL  
...  
for(i = first; i < n; ++i )  
    if(a[first+i] is desired item)  
        return &a[first+i];  
  
// if we drop through loop, then desired item was not found  
return NULL;
```

Serial Search Analysis

- What are the worst and average case running times for serial search?
- We must determine the O-notation for the number of operations required in search.
- Number of operations depends on n , the number of entries in the list.

Worst Case Time for Serial Search

- For an array of n elements, the worst case time for serial search requires n array accesses: $O(n)$.
- Consider cases where we must loop over all n records:
 - desired record appears in the last position of the array
 - desired record does not appear in the array at all

Average Case for Serial Search

Assumptions:

1. All keys are equally likely in a search
2. We always search for a key that is in the array

Example:

- We have an array of 10 records.
- If search for the first record, then it requires 1 array access; if the second, then 2 array accesses. *etc.*

The average of all these searches is:

$$(1+2+3+4+5+6+7+8+9+10)/10 = 5.5$$

Average Case Time for Serial Search

Generalize for array size n .

Expression for average-case running time:

$$(1+2+\dots+n)/n = n(n+1)/2n = (n+1)/2$$

Therefore, average case time complexity for serial search is $O(n)$.

Binary Search

- Perhaps we can do better than $O(n)$ in the average case?
- Assume that we are given an array of records that is sorted. For instance:
 - an array of records with integer keys sorted from smallest to largest (e.g., ID numbers), or
 - an array of records with string keys sorted in alphabetical order (e.g., names).

Binary Search Pseudocode

```
...
if(size == 0)
    found = false;
else {
    middle = index of approximate midpoint of array segment;
    if(target == a[middle])
        target has been found!
    else if(target < a[middle])
        search for target in area before midpoint;
    else
        search for target in area after midpoint;
}
...

```


Binary Search

Example: sorted array of integer keys. Target=7.

[0]	[1]	[2]	[3]	[4]	[5]	[6]
3	6	7	11	32	33	53

Binary Search

Example: sorted array of integer keys. Target=7.



Binary Search

Example: sorted array of integer keys. Target=7.


[0]	[1]	[2]	[3]	[4]	[5]	[6]
3	6	7	11	32	33	53



Is $7 = \text{midpoint key?}$ NO.


Binary Search

Example: sorted array of integer keys. Target=7.



Binary Search

Example: sorted array of integer keys. Target=7.



Search for the target in the area before midpoint.

Binary Search

Example: sorted array of integer keys. Target=7.

[0]	[1]	[2]	[3]	[4]	[5]	[6]
3	6	7	11	32	33	53



Find approximate midpoint

Binary Search

Example: sorted array of integer keys. Target=7.

[0]	[1]	[2]	[3]	[4]	[5]	[6]
3	6	7	11	32	33	53



Target = key of midpoint? NO.

Binary Search

Example: sorted array of integer keys. Target=7.

[0]	[1]	[2]	[3]	[4]	[5]	[6]
3	6	7	11	32	33	53



Target < key of midpoint? NO.

Binary Search

Example: sorted array of integer keys. Target=7.


[0]	[1]	[2]	[3]	[4]	[5]	[6]
3	6	7	11	32	33	53



Target > key of midpoint? YES.

Binary Search


Example: sorted array of integer keys. Target=7.



Search for the target in the area after midpoint.

Binary Search

Example: sorted array of integer keys. Target=7.



Binary Search Implementation


```
void search(const int a[ ], size_t first, size_t size, int target, bool& found, size_t& location)
{
    size_t middle;
    if(size == 0) found = false;
    else {
        middle = first + size/2;
        if(target == a[middle]){
            location = middle;
            found = true;
        }
        else if (target < a[middle])
            // target is less than middle, so search subarray before middle
            search(a, first, size/2, target, found, location);
        else
            // target is greater than middle, so search subarray after middle
            search(a, middle+1, (size-1)/2, target, found, location);
    }
}
```

Relation to Binary Search Tree

Array of previous example:

3	6	7	11	32	33	53
---	---	---	----	----	----	----

Corresponding complete binary search tree




Search for target = 7

Find midpoint:


3	6	7	11	32	33	53
---	---	---	----	----	----	----

Start at root:




Search for target = 7

Search left subarray:




Search left subtree:




Search for target = 7

Find approximate midpoint of subarray:




Visit root of subtree:




Search for target = 7

Search right subarray:



Search right subtree:



Binary Search: Analysis

- Worst case complexity?
- What is the maximum depth of recursive calls in binary search as function of n ?
- Each level in the recursion, we split the array in half (divide by two).
- Therefore maximum recursion depth is $\text{floor}(\log_2 n)$ and worst case = $O(\log_2 n)$.
- Average case is also = $O(\log_2 n)$.


Can we do better than $O(\log_2 n)$?

- Average and worst case of serial search = $O(n)$
- Average and worst case of binary search = $O(\log_2 n)$
- Can we do better than this?

YES. Use a hash table!

What is a Hash Table ?

- The simplest kind of hash table is an array of records.
- This example has 701 records.



[4]

Number 506643548

What is a Hash Table?

- Each record has a special field, called its key.
- In this example, the key is a long integer field called Number.

[0] [1] [2] [3]



[700]

...



[4]

What is a Hash Table?

- The number might be a person's identification number, and the rest of the record has information about the person.

[0] [1] [2] [3]




[700]

...

What is a Hash Table ?

- When a hash table is in use, some spots contain valid records, and other spots are "empty".



Open Address Hashing

- In order to insert a new record, the key must somehow be converted to an array index.
- The index is called the hash value of the key.



[0] [1] [2] [3] [4] [5] [700]



Inserting a New Record

- Typical way create a hash value:
(Number mod 701)



What is $(580625685 \% 701)$?

[0] [1] [2] [3] [4] [5] [700]



- Typical way to create a hash value:
(Number mod 701)

What is $(580625685 \% 701)$?



[0] [1] [2] [3] [4] [5] [700]



- The hash value is used for the location of the new record.



Inserting a New Record

- The hash value is used for the location of the new record.




Collisions

- Here is another new record to insert, with a hash value of 2.




**My hash
value is [2].**



Collisions

- This is called a collision, because there is already another valid record at [2].


When a collision occurs, move forward until you find an empty spot.



Collisions

- This is called a collision, because there is already another valid record at [2].


When a collision occurs, move forward until you find an empty spot.



Collisions

- This is called a collision, because there is already another valid record at [2].

When a collision occurs, move forward until you find an empty spot.




Collisions

- This is called a collision, because there is already another valid record at [2].

The new record goes in the empty spot.

[0] [1] [2] [3] [4] [5] [700]



Searching for a Key

- The data that's attached to a key can be found fairly quickly.

Number 701466868

[0] [1] [2] [3] [4] [5] [700]



- Calculate the hash value.
 - Check that location of the array for the key.

Number 701466868

My hash value is [2].

Not me.

[0] [1] [2] [3] [4] [5] [700]



Number 281942902



Number 233667136



Number 580625685



Number 506643548



Number 701466868



三



Number 155778322


- Keep moving forward until you find the key, or you reach an empty spot.

Number 701466868

My hash value is [2].

Not me.

[0] [1] [2] [3] [4] [5] [6] [700]




- Keep moving forward until you find the key, or you reach an empty spot.

Number 701466868

My hash value is [2].

Not me.

[0] [1] [2] [3] [4] [5] [6] [700]



- Keep moving forward until you find the key, or you reach an empty spot.

Number 701466868

My hash value is [2].

Yes!

[0] [1] [2] [3] [4] [5] [700]




- When the item is found, the information can be copied to the necessary location.



My hash value is [2].

Yes!

[0] [1] [2] [3] [4] [5] [700]




Deleting a Record

- Records may also be deleted from a hash table.




Deleting a Record

- Records may also be deleted from a hash table.
- But the location must not be left as an ordinary "empty spot" since that could interfere with searches.



Deleting a Record

- Records may also be deleted from a hash table.
- But the location must not be left as an ordinary "empty spot" since that could interfere with searches.
- The location must be marked in some special way so that a search can tell that the spot used to have something in it.



Hashing

- Hash tables store a collection of records with keys.
- The location of a record depends on the hash value of the record's key.
- Open address hashing:
 - When a collision occurs, the next available location is used.
 - Searching for a particular key is generally quick.
 - When an item is deleted, the location must be marked in a special way, so that the searches know that the spot used to be used.
- See text for implementation.

Open Address Hashing

- To reduce collisions...
 - Use table CAPACITY = prime number of form $4k+3$
 - Hashing functions:
 - Division hash function: key \% CAPACITY
 - Mid-square function: $(\text{key} * \text{key}) \% \text{CAPACITY}$
 - Multiplicative hash function: key is multiplied by positive constant less than one. Hash function returns first few digits of fractional result.

Clustering

- In the hash method described, when the insertion encounters a collision, we move forward in the table until a vacant spot is found. This is called *linear probing*.
- *Problem:* when several different keys are hashed to the same location, adjacent spots in the table will be filled. This leads to the problem of *clustering*.
- As the table approaches its capacity, these clusters tend to merge. This causes insertion to take a long time (due to linear probing to find vacant spot).

Double Hashing

- One common technique to avoid cluster is called *double hashing*.
- Let's call the original hash function *hash1*
- Define a second hash function *hash2*

Double hashing algorithm:

1. *When an item is inserted, use hash1(key) to determine insertion location i in array as before.*
2. *If collision occurs, use hash2(key) to determine how far to move forward in the array looking for a vacant spot:*

$$\text{next location} = (i + \text{hash2}(\text{key})) \% \text{CAPACITY}$$

Double Hashing


- Clustering tends to be reduced, because `hash2()` has different values for keys that initially map to the same initial location via `hash1()`.
- This is in contrast to hashing with *linear probing*.
- Both methods are *open address hashing*, because the methods take the next open spot in the array.
- In linear probing
$$\text{hash2(key)} = (i+1)\% \text{CAPACITY}$$
- In double hashing `hash2()` can be a general function of the form
 - $\text{hash2(key)} = (l+f(key))\% \text{CAPACITY}$

Chained Hashing

- In open address hashing, a collision is handled by probing the array for the next vacant spot.
- When the array is full, no new items can be added.
- We can solve this by resizing the table.
- Alternative: chained hashing.

Chained Hashing

- In chained hashing, each location in the hash table contains a list of records whose keys map to that location:



Time Analysis of Hashing

- Worst case: every key gets hashed to same array index! $O(n)$ search!!
- Luckily, average case is more promising.
- First we define a fraction called the hash table *load factor*:

$$\alpha = \frac{\text{number of occupied table locations}}{\text{size of table's array}}$$

Average Search Times

For open addressing with linear probing, average number of table elements examined in a successful search is approximately:

$$\frac{1}{2} (1 + \frac{1}{1-\alpha})$$

Double hashing: $-\ln(1-\alpha)/\alpha$

Chained hashing: $1+\alpha/2$

Average number of table elements examined during successful search

Load factor(α)	Open addressing, linear probing $\frac{1}{2} (1+1/(1-\alpha))$	Open addressing double hashing $-\ln(1-\alpha)/\alpha$	Chained hashing $1+\alpha/2$
0.5	1.50	1.39	1.25
0.6	1.75	1.53	1.30
0.7	2.17	1.72	1.35
0.8	3.00	2.01	1.40
0.9	5.50	2.56	1.45
1.0	Not applicable	Not applicable	1.50
2.0	Not applicable	Not applicable	2.00
3.0	Not applicable	Not applicable	2.50

Summary

- Serial search: average case $O(n)$
- Binary search: average case $O(\log_2 n)$
- Hashing
 - Open address hashing
 - Linear probing
 - Double hashing
 - Chained hashing
 - Average number of elements examined is function of load factor α .